



Chapter 1

Feed-Forward Neural Nets


1.1 Perceptrons

- Find a line that separates red and blue

x_1	x_2	a
0.36	0.24	1
0.82	0.09	1
-0.48	1.00	0
\vdots	\vdots	\vdots





- $\Phi = \{w, b\}$
 where $w = [w_1, \dots, w_m]$, $b \in \mathbb{R}$
- $f_{\Phi}(\mathbf{x}) = \begin{cases} 1 & \text{if } b + w \cdot \mathbf{x} > 0 \\ 0 & \text{otherwise} \end{cases}$



- Algorithm
 1. set $b = 0, w_i = 0$
 2. until the weights do not change
 - (a) for each training example \mathbf{x}^k with answer a^k
 - i. if $a^k - f(\mathbf{x}^k) = 0$, then continue
 - ii. else $w_i \leftarrow w_i + (a^k - f(\mathbf{x}^k)) x_i^k$

Multiclass Decision Problem



1.2 Cross-entropy Loss Functions for Nerual Nets

Softmax

- Definition:

$$\sigma(\mathbf{x})_j = \frac{e^{x_j}}{\sum_i e^{x_i}}$$

$$\sigma(x_1, \dots, x_m) = \left(\frac{e^{x_1}}{\sum_i e^{x_i}}, \dots, \frac{e^{x_m}}{\sum_i e^{x_i}} \right)$$

Cross-entropy

- Definition:

$$H(p, q) = \sum_i p_i \log q_i$$

Cross-entropy loss

- Definition:

$$X(\Phi, \mathbf{x}) = -\ln p_{\Phi}(a_{\mathbf{x}})$$

where $p_{\Phi}(a_x)$ is the probability assigned to \mathbf{x} 's label

1.3 Derivatives and Stochastic Gradient Descent

- Equations:

$$X(\Phi, \mathbf{x}) = -\ln p(a)$$

$$p(a) = \sigma_a(\mathbf{l}) = \frac{e^{l_a}}{\sum_i e^{l_i}}$$

$$l_j = b_j + \mathbf{x} \cdot \mathbf{w}_j$$

where $\mathbf{l} = (l_1, \dots, l_m)$

- Gradients:

$$\begin{aligned}\frac{\partial X(\Phi)}{\partial b_j} &= \frac{\partial l_j}{\partial b_j} \frac{\partial X(\Phi)}{\partial l_j} \\ &= \begin{cases} -(1 - p_j) & \text{if } a = j \\ p_j & \text{otherwise} \end{cases}\end{aligned}$$

$$\begin{aligned}\frac{\partial X(\Phi)}{\partial w_{i,j}} &= \frac{\partial l_j}{\partial w_{i,j}} \frac{\partial X(\Phi)}{\partial l_j} \\ &= \begin{cases} -(1 - p_j)x_i & \text{if } a = j \\ p_j x_i & \text{otherwise} \end{cases}\end{aligned}$$

- Update: For learning rate \mathcal{L}

$$b_j \leftarrow b_j - \mathcal{L} \frac{\partial X(\Phi)}{\partial b_j}$$

$$w_{i,j} \leftarrow w_{i,j} - \mathcal{L} \frac{\partial X(\Phi)}{\partial w_{i,j}}$$

1.4 Writing Our Program

- Data Normalization
 - $-1 \leq x_i \leq 1$
- Learning Rate
 - 0.0001 in MNIST
- Weights and Bias
 - $[-0.1, 0.1]$

1.5 Matrix Representation of Neural Nets

- Forward Propagation:

$$\mathbf{L} = \mathbf{XW} + \mathbf{B}$$

where \mathbf{X} is the input matrix

- Loss:

$$\Pr(A(\mathbf{x})) = \sigma(\mathbf{xW} + \mathbf{b})$$

$$L(\mathbf{x}) = -\log(\Pr(A(\mathbf{x}) = a))$$

- Update:

$$\Delta \mathbf{W} = -\mathcal{L} \mathbf{X}^T \nabla_{\mathbf{L}} X(\Phi)$$

1.6 Data Independence

- iid assumption