

Chapter 4

Word Embeddings and Recurrent NNs

4.1 Word Embeddings for Language Models

- Variable $E_{1,n} = (E_1 \dots E_n)$

- $E_i = \text{a word}$

- Value $e_{1,n}$

- E.g. $e_{1,6} = (\text{We live in a small world})$
 $e_1 = \text{We}, e_2 = \text{live}, \dots$

- Probability Model

$$p(\text{We live in a small World}) = p(\text{We})p(\text{live}|\text{We})p(\text{in}|\text{We live}) \dots$$

- Generally

$$p(E_{1,n} = e_{1,n}) = \prod_{j=1}^{j=n} p(E_j = e_j | E_{1,j-1} = e_{1,j-1})$$

- Tokenization
 - Breaking the strings into a sequence of words
 - Vocabulary set V
 - Punctuation(final period .) is a word
 - *UNK* is unknown words
 - E.g. "132,423" in the sentence "The population of Providence is 132,423."

- Penn Treebank(PTB)

- news articles from the Wall Street Journal.
- Files
 - ptb.train.txt, ptb.test.txt, ptb.valid.txt
- Train data
 - 929589 words
- Test data
 - 82430 words
- Vocabulary set V
 - 10000 words

aer banknote berlitz calloway
centrust cluett fromstein gitano
guterman hydro-quebec ipo kia
memotec mlx nahb punts rake
regatta rubens sim snack-food
ssangyong swapo wachter
pierre <unk> N years old will
join the board as a nonexecutive
director nov. N
mr. <unk> is chairman of <unk>
n.v. the dutch publishing group

- Probability Model

- Bigram model

$$p(E_{1,n} = e_{1,n}) = p(E_1 = e_1) \prod_{j=2}^{j=n} p(E_j = e_j | E_{j-1} = e_{j-1})$$

- Sentence padding

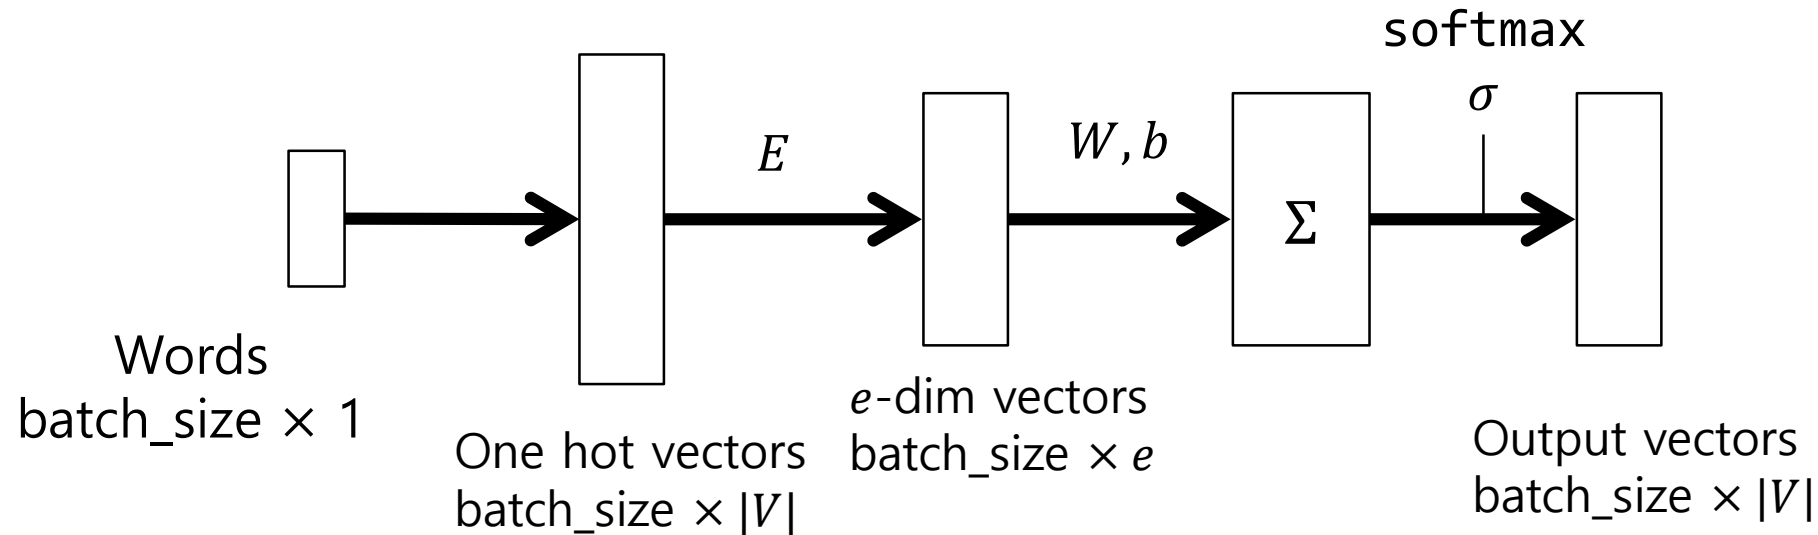
- Put a imaginary word "STOP" at the beginning of every sentence

$$p(E_{1,n} = e_{1,n}) = \prod_{j=1}^{j=n} p(E_j = e_j | E_{j-1} = e_{j-1})$$

- Neural Network

- Word embedding

- word \rightarrow index \rightarrow one hot vector $\rightarrow e$ -dimensional vector
 - $e = 20, 100$ or 1000



- Words embedding
 - Cosine similarity
 - $\cos(\mathbf{x}, \mathbf{y}) = \frac{\mathbf{x} \cdot \mathbf{y}}{\|\mathbf{x}\| \|\mathbf{y}\|}$
 - Cosine of angle between \mathbf{x} and \mathbf{y}

Word Numbers	Word	Largest Cosine Similarity	Most Similar
0	under		
1	above	0.362	0
2	the	-0.160	0
3	a	0.127	2
4	recalls	0.479	1
5	says	0.553	4
6	rules	-0.066	4
7	laws	0.523	6
8	computer	0.249	2
9	machine	0.333	8

4.2 Building Feed-Forward Language Models

```
inpt = tf.placeholder(tf.int32, shape=[batchSz])
answr = tf.placeholder(tf.int32, shape=[batchSz])
E = tf.Variable(tf.random_normal([vocabSz, embedSz], stddev = 0.1))
embed = tf.nn.embedding_lookup(E, inpt)
```

```
xEnt = tf.nn.sparse_softmax_cross_entropy_with_logits(logits=logits, labels=answr)
loss = tf.reduce_sum(xEnt)
```


4.3 Improving Feed-Forward Language Models

1. Adding layer

2. Trigram model

```
embed2 = tf.nn.embedding_lookup(E, inpt2)  
both = tf.concat([embed, embed2], 1)
```

4.4 Overfitting

- Regularization
 - Early stopping
 - Dropout
 - L2 regularization

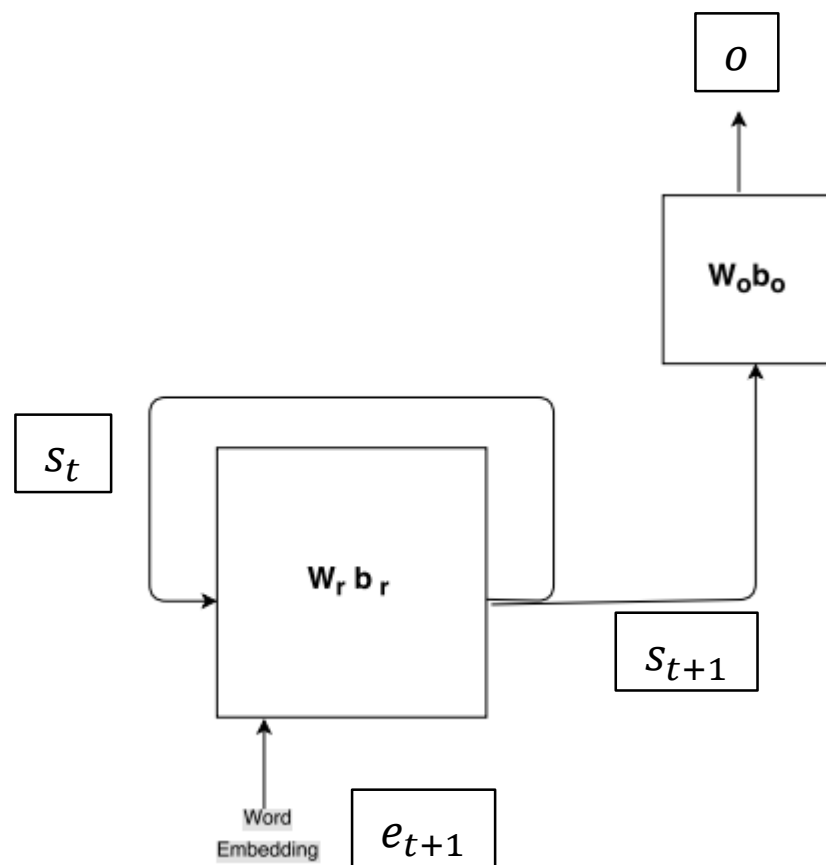
Epoch	1	2	3	4	5	6	7	10	15	20	30
Train	197	122	100	87	78	72	67	56	45	41	35
Dev	172	152	145	143	143	143	145	149	159	169	182

Figure 4.3: Overfitting in a language model

Epoch	1	2	3	4	5	6	7	10	15	20	30
Dropout	213	182	166	155	150	144	139	131	122	118	114
L2 Reg	180	163	155	148	144	140	137	130	123	118	112

Figure 4.4: Language-model perplexity when using regularization

4.5 Recurrent Networks




RNN(recurrent neural network)

$$\begin{aligned}s_0 &= 0 \\ s_{t+1} &= \rho((e_{t+1} \cdot s_t)W_r + b_r) \\ o &= \sigma(s_{t+1}W_o + b_o)\end{aligned}$$

ρ : relu
 σ : cross-entropy
 \cdot : concatenation

It is a small world but I like it that way



batch 1

batch 2

STOP	It	is	a	small	world
but	I	like	it	that	way

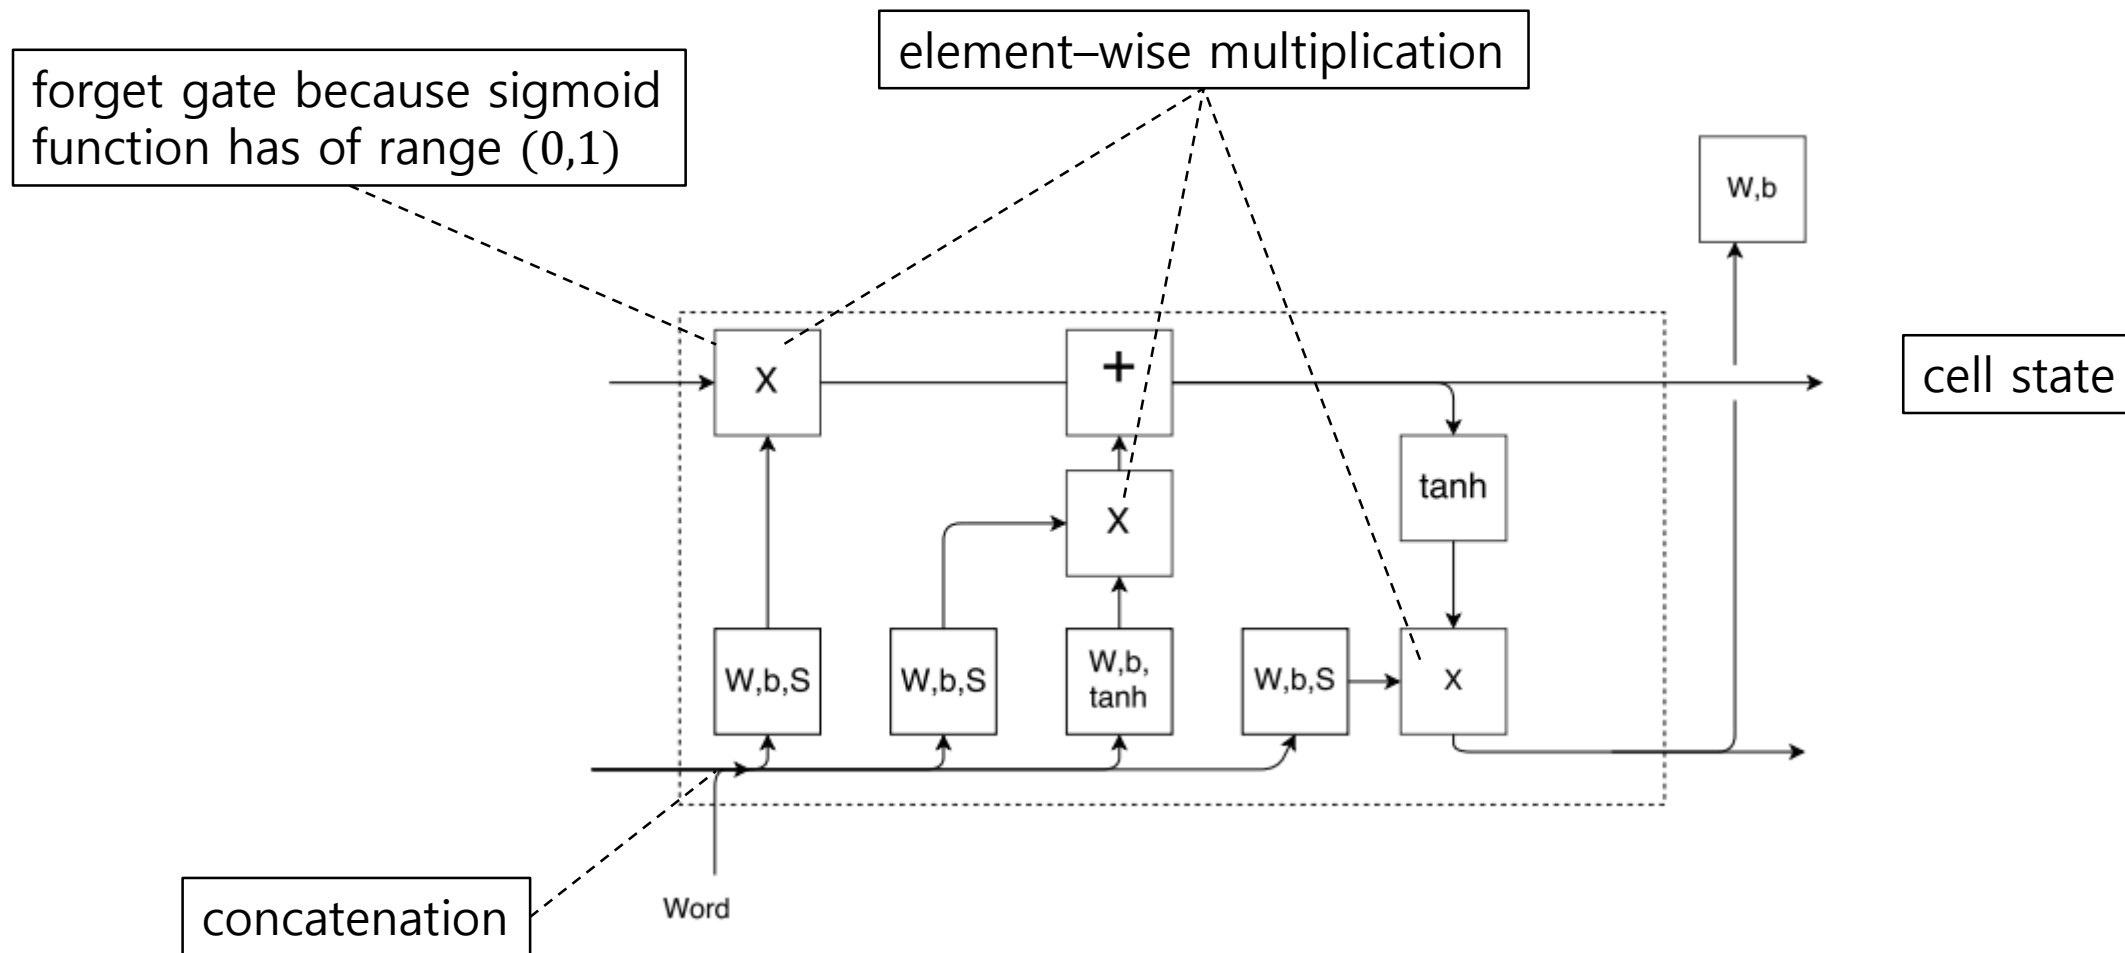
STOP	It	is
but	I	like

a	small	world
it	that	way

Batch size = 2
Window size = 3

```
rnn= tf.contrib.rnn.BasicRNNCell(rnnSz)
initialState = rnn.zero_state(batchSz, tf.float32)
outputs, nextState = tf.nn.dynamic_rnn(rnn, embeddings,
                                       initial_state=initialState)
```

4.6 Long Short-Term Memory



- TF program

- replace

- `tf.contrib.rnn.BasicRNNCell(rnnSz)`

- with

- `tf.contrib.rnn.LSTMCell(rnnSz)`